

January 20, 2018

1 Resizing image Using Interpolation

1.1 Importing images

```
In [5]: img1 = imread('./blur.jpg');
```

```
In [21]: img2 = imread('./cameraman.png');
```

```
In [8]: img3 = imread('./shapes.gif');
```


1.2 Using Nearest Neighbours

In nearest neighbour the closest pixel values is copied to the pixel which results in a more pixelated image which can be used accordingly in particular types of art

```
In [12]: new_img1 = RESIZENN(img1,5);  
         imshowpair(img1,new_img1,'montage');
```

Warning: Image is too big to fit on screen; displaying at 8%

```
> In images.internal.initSize (line 71)  
   In imshow (line 328)  
   In imshowpair (line 126)
```



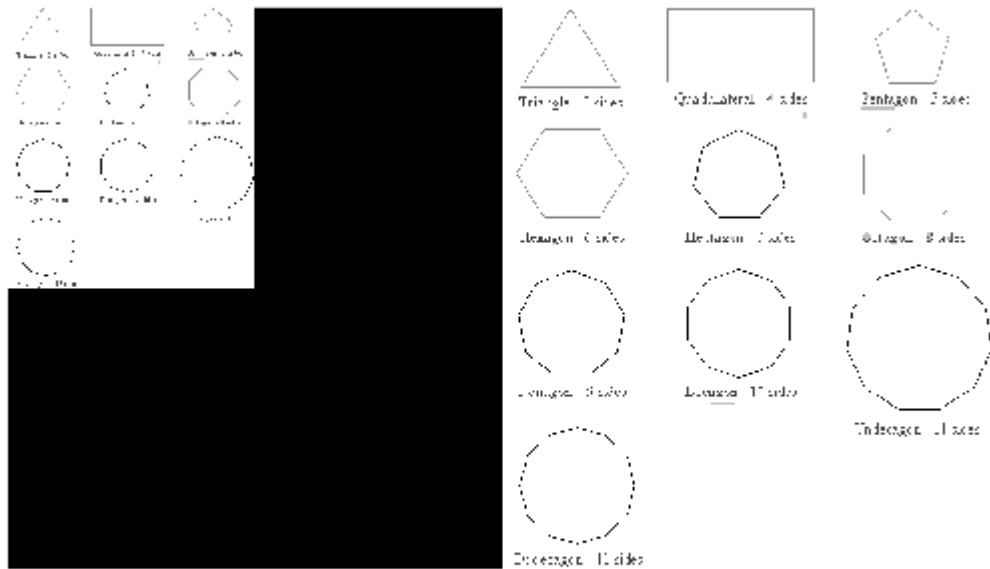
```
In [23]: new_img2= RESIZENN(img2,2);  
         imshowpair(img2,new_img2,'montage');
```



```
In [22]: new_img3= RESIZENN(img3,2);  
         imshowpair(img3,new_img3,'montage');
```

Warning: Image is too big to fit on screen; displaying at 67%


```
> In images.internal.initSize (line 71)  
   In imshow (line 328)  
   In imshowpair (line 126)
```




2 Using Bilinear Interpolation

Bilinear interpolation results in a smoother image due to rowwise linear interpolation i.e the contribution of the closer pixel is more than the farther one but still both contribute this results in a smoother picture.


```
In [25]: red = [0.4 1 0; 0 1 1; 0.8 0.8 0.1 ];
         green = [1 0.5 0; 0 0 1; 0.5 0.3 0.7];
         blue = [0.0 0.0 0; 1.0 1.0 0.2; 1.0 0.4 1];
         img = cat(3, red, green, blue);
         image(img);
```



```
In [32]: new_img = RESIZEBL(img,2);  
         image(new_img)
```



```
In [33]: new_img = RESIZEBL(img,3);  
         image(new_img)
```



```
In [34]: new_img1 = RESIZEBL(img1,5);
         imshowpair(img1,new_img1,'montage');
```

Warning: Image is too big to fit on screen; displaying at 8%

```
> In images.internal.initSize (line 71)
   In imshow (line 328)
   In imshowpair (line 126)
```



```
In [35]: new_img2= RESIZEBL(img2,2);  
         imshowpair(img2,new_img2,'montage');
```




2.0.1 The Difference between the two types of interpolations

The major difference between the two is that the images resized by nearest neighbour are more pixelated and those done by bilinear are not. Nearest neighbour works equal or even better with images with higher number of straight edges.

We have a lot of other alternatives for image resizing like bicubic interpolation which uses the cubic equations for finding the pixel values.

In []:

q2

January 20, 2018

1 Applying and implementing convolution


```
In [3]: img = imread('./cameraman.png');  
        imshow(img)  
        size(img);
```



```
In [6]: M = [1 2 1; 0 0 0; -1 -2 -1];  
        M = double(M);
```

2 Detecting Horizontal Edges


```
In [9]: new_img = conv2(M,img);  
        new_img = mat2gray(new_img);  
        imshow(new_img)
```



3 Detecting Vertical Edges

By taking the transpose we can find Vertical images

```
In [10]: new_img = conv2(M',img);  
         new_img = mat2gray(new_img);  
         imshow(new_img)
```



Originally we can't see much difference but after converting it to gray we can see the edges detected

In []:

q3

January 20, 2018

1 Outsize after convolution

We are given - Image size = Width, Height, Channels - No. of filters = N - filter size = F,F,Channels - Stride = S - padding = Z

Output Size = $N(\text{Width} - F + 2P)/S + 1, N(\text{Height} - F + 2P)/S + 1, \text{Channels}$

2 No. of Multiplications:

Multiplications in 1 Operation = $\frac{FF}{\text{Output_width} \times \text{Output_Height} \times \text{Channels}}$ Addition in 1 Operation = 1 Total Operations =

$\frac{\text{Total Multiplication} = (FF)}{\text{Output_width} \times \text{Output_Height} \times \text{Channels}}$ Total Addition = 1 * Out-

Reference

In []:

q4

January 20, 2018

1 Audio Convolution

```
In [10]: audio_file = audiorecorder(44100,24,1,1)
```

```
audio_file =
```

```
audiorecorder with properties:
```

```
    SampleRate: 44100
    BitsPerSample: 24
    NumberOfChannels: 1
    DeviceID: 1
    CurrentSample: 1
    TotalSamples: 0
    Running: 'off'
    StartFcn: []
    StopFcn: []
    TimerFcn: []
    TimerPeriod: 0.0500
    Tag: ''
    UserData: []
    Type: 'audiorecorder'
```

```
In [11]: disp('Start speaking.')
         recordblocking(audio_file, 5);
         disp('End of Recording.');
```


```
Start speaking.
End of Recording.
```

```
In [40]: play(audio_file);
```


```
In [26]: y = getaudiodata(audio_file);
         size_y = size(y,1)
         plot(y);
```

```
size_y =
```


```
220500
```




```
In [30]: y24 = y((1:floor(44.1/24):size_y)');  
         P24 = audioplayer(y24,24000);  
         plot(y24)  
         play(P24)
```




```
In [31]: y16 = y((1:floor(44.1/16):size_y)');  
         P16 = audioplayer(y24,16000);  
         plot(y16)  
         play(P16)
```


```
In [32]: y8 = y((1:floor(44.1/8):size_y)');  
         P8 = audioplayer(y24,8000);  
         plot(y8)  
         play(P8)
```




```
In [33]: y4 = y((1:floor(44.1/4):size_y)');  
         P4 = audioplayer(y24,4000);  
         plot(y4)  
         play(P4)
```




```
In [34]: church_bullet = audioread('church_bullet.wav');  
         church = audioplayer(church_bullet,44100);  
         play(church);  
         plot(church_bullet);
```




```
In [38]: church_y = conv(y,church_bullet(:,1));  
         church_conv_audio = audioplayer(church_y,44100);  
         plot(church_y)  
         play(church_conv_audio)
```



```
In [39]: r1_bullet = audioread('./r1_omni.wav');  
         r1_y = conv(y,r1_bullet(:,1));  
         r1_conv_audio = audioplayer(r1_y,44100);  
         plot(r1_y)  
         play(r1_conv_audio)
```



```
In [41]: tr_bullet = audioread('./terrrys_typing_omni.wav');  
         tr_y = conv(y,tr_bullet(:,1));  
         tr_conv_audio = audioplayer(tr_y,44100);  
         plot(tr_y)  
         play(tr_conv_audio)
```



In []:

q5

January 20, 2018

1 2 D Convolution and edge detection

Let assume the below 3x3 matrix

```
In [22]: M = [1,2,1;0,0,0;-1,-2,-1]
```

M =

```
1    2    1
0    0    0
-1   -2   -1
```

We can see after convolution, we get a red line.

```
In [10]: img = imread('./sample_inp.png');
         new_img = myconv(M,img);
```

```
In [11]: imshow(new_img);
```




1.0.1 Using the M we first detect many vertical edges

```
In [24]: blur = imread('./blur.jpg');  
        ver_blur = myconv(M,blur);  
        imshow(ver_blur*5)
```

Warning: Image is too big to fit on screen; displaying at 67%

```
> In images.internal.initSize (line 71)  
   In imshow (line 328)
```




1.0.2 Using the M' we detect many horizontal edges

```
In [25]: hor_blur = myconv(M',blur);  
        imshow(hor_blur*5)
```

Warning: Image is too big to fit on screen; displaying at 67%


```
> In images.internal.initSize (line 71)  
   In imshow (line 328)
```



```
In [30]: fin_blur = ver_blur + hor_blur;  
         imshow(fin_blur.*5)
```

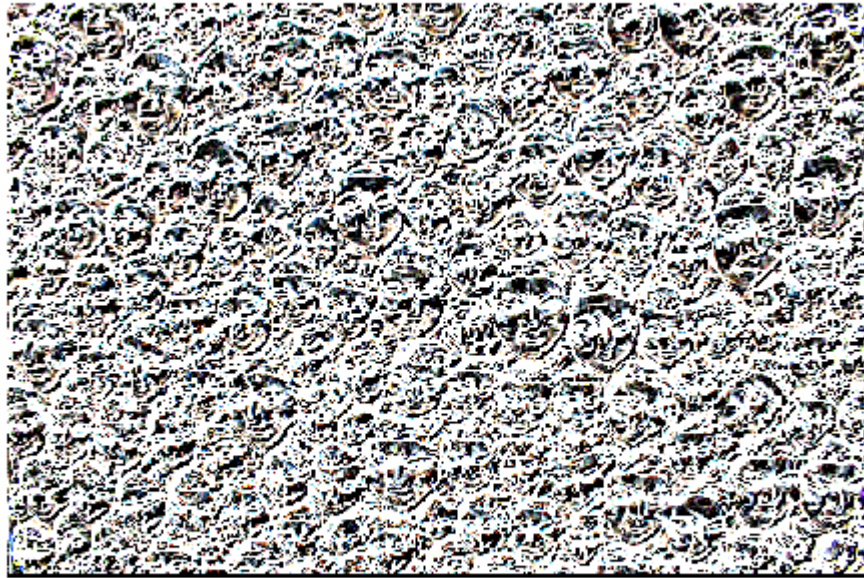
Warning: Image is too big to fit on screen; displaying at 67%

```
> In images.internal.initSize (line 71)  
   In imshow (line 328)
```



Hence we can even do this on other images'

```
In [33]: img = imread('./Faces.jpg');  
         ver_img = myconv(M,img);  
         hor_img = myconv(M',img);  
         fin_img = ver_img + hor_img;  
         imshow(fin_img.*5)
```

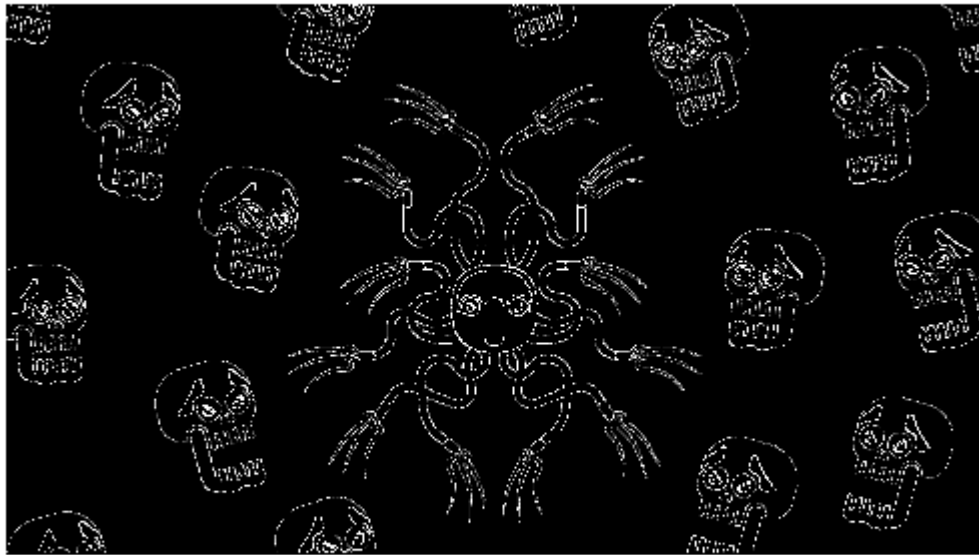



```
In [34]: img = imread('./War_on_drugs.png');  
         ver_img = myconv(M,img);  
         hor_img = myconv(M',img);  
         fin_img = ver_img + hor_img;  
         imshow(fin_img.*5)
```

Warning: Image is too big to fit on screen; displaying at 67%

> In images.internal.initSize (line 71)

In imshow (line 328)




1.0.3 We get better edge detection after using some kind of blur

```
In [36]: fin_img = conv2([1,1,1;1,1,1;1,1,1]./9,fin_img);  
         imshow(fin_img)
```

Warning: Image is too big to fit on screen; displaying at 67%

```
> In images.internal.initSize (line 71)  
   In imshow (line 328)
```




In []:

q6

January 20, 2018

1 Sub Image Detection using Cross Correlation

```
In [2]: img = imread('./Faces.jpg');  
        image(img)  
        img = rgb2gray(img);
```



```
In [3]: sub_img = imread('./F1.jpg');  
        image(sub_img)  
        sub_img = rgb2gray(sub_img);
```



```
In [4]: new_img = normxcorr2(sub_img(:,:,1),img(:,:,1));
```

```
In [5]: [val,p] = max(new_img(:));
        x = floor(p/size(new_img,1))
        y = p - x*size(new_img,1)
```


```
x =
```

```
453
```


```
y =
```

```
132
```

```
In [6]: imshow(new_img);
        hold on;
        % Then, from the help:
        rectangle('Position',[x - size(sub_img,1)/2,y - size(sub_img,2)/2,size(sub_img,1),size(sub_img,2)],...
                  'Curvature',[0.8,0.4],...
                  'LineWidth',2,'LineStyle','--','EdgeColor','r')
```

```
In [8]: imshow(img);
        hold on;
        % Then, from the help:
        rectangle('Position',[x - size(sub_img,1),y - size(sub_img,2),size(sub_img,1),size(sub_img,2)],
                  'Curvature',[0.8,0.4],...
                  'LineWidth',2,'LineStyle','--','EdgeColor','r')
```



1.1 Report

- We can see various white spots in the final correlation image
- These dots represent the highly similar sections with the corresponding sub_image
- Hence we can say we detected the subimage at these locations in the original image

In []:

q7

January 20, 2018

1 Finding the 1-D filter

```
In [3]: inp = [12, 20, 3, 10, 22, 19, 23, 16, 0, 21, 23, 16, 18]
        inp_size = size(inp,2)
```

```
inp =
```

```
    12    20     3    10    22    19    23    16     0    21    23    16    18
```

```
inp_size =
```

```
    13
```

```
In [4]: out = [75, 52, 33, 97, 251, 211, 63, 65]
        out_size = size(out,2)
```

```
out =
```

```
    75    52    33    97   251   211    63    65
```

```
out_size =
```

```
     8
```

1.0.1 Method:

- Using Linear regressio we will try to find the filter
- Hence we are assuming random values in beginning
- We will calculate the pred_out
- We will find the loss and gradient

1.1 Using Normalize equations

```
In [7]: % Assuming stride = 1
        % Assuming padding = 0
        filter_size = (size(inp,2) - size(out,2) + 1);
```

```
In [8]: inp_pred = zeros(filter_size,out_size);
        for i = 1:inp_size - filter_size + 1
            inp_pred(:,i) = inp(i:i+filter_size-1);
        end
        inp_pred
```

```
inp_pred =

    12    20     3    10    22    19    23    16
    20     3    10    22    19    23    16     0
     3    10    22    19    23    16     0    21
    10    22    19    23    16     0    21    23
    22    19    23    16     0    21    23    16
    19    23    16     0    21    23    16    18

In [9]: filter = out*pinv(inp_pred)
filter =

    5.0000    4.0000    4.0000   -3.0000   -3.0000    1.0000
```

1.2 Using Gradient Descent

```
In [10]: epoch = 1000;
         lr = 0.001;
         FILTER = rand(1,filter_size).*10

FILTER =

    8.1472    9.0579    1.2699    9.1338    6.3236    0.9754

In [18]: for ind = 1:epoch

         out_pred = FILTER*inp_pred;
         loss = sum((out - out_pred).^2)/(2*out_size);
         gradient = (out_pred - out)*inp_pred'/out_size;
         FILTER = FILTER - lr.*gradient;

         end

In [16]: FILTER = round(FILTER)
         out_pred = FILTER*inp_pred
         out

FILTER =

     5     4     4    -3    -3     1

out_pred =

    75    52    33    97   251   211    63    65

out =

    75    52    33    97   251   211    63    65

In [17]: FILTER = flip(FILTER)
FILTER =

     1    -3    -3     4     4     5

In [ ]:
```